

Time series modeling with Bayesian Dynamic Generalized Additive Models


Nicholas Clark

School of Veterinary Science, University of Queensland, Australia


Wednesday 13th December, 2023




GAMs use splines ($f(x)$) ...



...made of weighted basis functions



Penalize $f''(x)$ to learn weights



Easy to fit in


$$\mathbb{E}(\mathbf{Y}_t | \mathbf{X}_t) = g^{-1}\left(\alpha + \sum_{j=1}^J f(x_{jt})\right)$$


Where:

g^{-1} is the inverse of the link function


α is the intercept

$f(x)$ are potentially nonlinear functions of the J predictors







Basis: GP



Basis: Tensor product







Basis: TPRS



Basis: CRS (shrink)

What's the catch?




A spline of time


```
library(mgcv)
model <- gam(y ~ s(time, k = 20, bs = 'bs', m = 2),
             data = data,
             family = gaussian())
```

A B-spline (`bs = 'bs'`) with `m = 2` sets the penalty on the second derivative


Hindcasts ☺




Basis functions \Rightarrow local knowledge




Basis functions \Rightarrow local knowledge




Forecasts



We need *global* knowledge



We need *global* knowledge



Dynamic GAMs

$$\mathbb{E}(\mathbf{Y}_t | \mathbf{X}_t) = g^{-1}(\alpha + \sum_{j=1}^J f(x_{jt}) + z_t)$$

Where:

g^{-1} is the inverse of the link function

α is the intercept

$f(x)$ are potentially nonlinear functions of the J predictors

z_t is a ***latent dynamic process***

Modelling with the mvgam

Bayesian framework to fit Dynamic GLMs and Dynamic GAMs

Hierarchical intercepts, slopes and smooths

Latent dynamic processes

State-Space models with measurement error

Built off the mgcv  to construct penalized smoothing splines

Familiar  formula interface

Uni- or multivariate series from a range of response distributions


Uses Stan for ADVI, Laplace or full Hamiltonian Monte Carlo

We can fit models that include random effects, nonlinear effects, time-varying effects and complex multidimensional smooth functions. All these effects can operate *on both process and observation models*


Can incorporate unobserved temporal dynamics; no need to regress the outcome on past values or resort to transformations

What kinds of dynamic processes are available in the `mvgam`  ?


Piecewise linear...




...or logistic with upper saturation




RW or ARMA($p = 1-3$, $q = 0-1$)




Gaussian Process...




...where length scale \Rightarrow *memory*



VAR1 \Rightarrow Granger causality



Factors \Rightarrow induced correlations



Example of the interface

```
model ← mvgam(  
  formula = y ~  
    s(series, bs = 're') +  
    s(x0, series, bs = 're') +  
    x1 +  
    gp(x2) +  
    te(x3, x4, bs = c('cr', 'tp')),  
  data = data,  
  family = poisson(),  
  trend_model = AR(p = 1, ma = TRUE, cor = TRUE),  
  algorithm = 'sampling',  
  burnin = 500,  
  samples = 500,  
  chains = 4,  
  parallel = TRUE)
```

Produce all Stan code and objects


```
code(model)
```

```
## // Stan model code generated by package mvgam
## functions {
##   /* Spectral density function of a Gaussian process
##   * with squared exponential covariance kernel
##   * Args:
##   *   l_gp: numeric eigenvalues of an SPD GP
##   *   alpha_gp: marginal SD parameter
##   *   rho_gp: length-scale parameter
##   * Returns:
##   *   numeric values of the GP function evaluated at l_gp
##   */
##   vector spd_cov_exp_quad(data vector l_gp, real alpha_gp, real rho_gp) {
##     int NB = size(l_gp);
##     vector[NB] out;
##     real constant = square(alpha_gp) * (sqrt(2 * pi()) * rho_gp);
##     real neg_half_lscales2 = -0.5 * square(rho_gp);
##     for (m in 1 : NB) {
##       out[m] = constant * exp(neg_half_lscales2 * square(l_gp[m]));
##     }
##     return out;
##   }
```

Workflow in `mvgam`

Fit models with splines, GPs, and multivariate dynamic processes to sets of time series; use informative priors for effective regularization

Use posterior predictive checks and Randomized Quantile residuals to assess model failures

Use `marginaleffects`  to generate interpretable (and reportable) model predictions

Produce probabilistic forecasts

Evaluate forecasts from using proper scoring rules

More resources

Cheatsheet ⇒ [Overview of `mvgam`](#)

Vignette ⇒ [Introduction to the package](#)

Vignette ⇒ [Shared latent process models](#)

Vignette ⇒ [Time-varying effects](#)

Vignette ⇒ [Multivariate State-Space models](#)

Motivating publication ⇒ Clark & Wells 2023 [*Methods in Ecology and Evolution*](#)

Relevant links

[mvgam](#)  website

 [nicholasjclark](#)

 [slides for this talk](#)

 [personal website](#)

 n.clark@uq.edu.au